

Title: Bringing Elton and Grinnell together: a quantitative framework to represent the biogeography of ecological interaction networks

Authors: Dominique Gravel^{1,2,*}, Benjamin Baiser³, Jennifer A. Dunne⁴, Jens-Peter Kopelke⁵, Neo Martinez⁶, Tommy Nyman⁷, Timothée Poisot^{2,8}, Spencer A. Wood⁹, Daniel B. Stouffer¹⁰, Jason Tylianakis^{10,11}, Tomas Roslin¹²,

1: Canada Research Chair in Integrative Ecology. Département de biologie, Université de Sherbrooke, 2500 Boulevard l'Université, Sherbrooke (Québec). J1K 2R1

2: Québec Centre for Biodiversity Sciences

3: University of Florida, Department of Wildlife Ecology and Conservation, 110 Newins-Ziegler Hall, PO Box 110430, Gainesville, Fl. 32611-0430

4:

5: Senckenberg Research Institute and Natural History Museum, Senckenberganlage 25, D-60325 Frankfurt am Main, Germany

6:

7:

8:

9: University of Washington, School of Environmental and Forest Sciences, Box 352100, Seattle, WA 98195, USA

10: University of Canterbury at Christchurch, School of Biological Sciences

11:

12: Department of Ecology, Swedish University of Agricultural Sciences, Box 7044, 750 07 Uppsala, Sweden

Keywords: networks, spatial ecology, co-occurrence, probability of interaction

Words in the abstract:

Words in the main text:

Words in the legends:

Figures:

Tables:

References:

Abstract

Biogeography has historically focused on the spatial distribution and abundance of species, neglecting variation in the way species interact with each other. Models of species distribution and of interactions have historically taken different paths. Here, we plea for an integrated approach, adopting the view that community structure is best represented as an ecological network of interactions. We outline a conceptual approach, suggesting that the ecological niche may be redefined to encompass the effect of the environment on species distribution (the Grinnellian dimension of the niche) and on the ecological interactions among them (the Eltonian dimension). Starting from this novel concept, we develop a quantitative theory to explain turnover of interactions in space and time – i.e. a novel approach to interaction distribution modelling. We apply this novel framework to a large data set of host–parasite interactions across Europe and find that two aspects of the environment (temperature and precipitation) leave a strong imprint on species co-occurrence, but not on the probability of local interactions. Even where species co-occur, interaction proves a stochastic rather than deterministic process, adding to variation in realized network structure. We also find that a large majority of species pairs are never found together, thus precluding any inferences regarding their probability to interact. Our framework provides a first conceptual framework to explain the variation of network structure at large spatial scales and opens new perspectives at the frontier between species distribution modelling and community ecology.

Introduction

Community ecology is defined in most textbooks as *the study of the interactions that determine the distribution and abundance of organisms* (?). Despite a general consensus on this definition (Scheiner2007), research on variation in community structure has mostly focused on the spatial and temporal turnover of species composition (?), neglecting variation in the way species interact with each other (?). Given this omission, it is perhaps not surprising that biogeographers are still struggling to establish whether interactions actually impact the distribution of species at large spatial scales(??), even though there is enough evidence from community ecology that interactions are critical for determining species dynamics, resource use and responses to the environment at finer scales. Recent attempts at accounting for interactions in species distribution models (??) have brought some methodological advances. Yet, these techniques are still based on a ‘species-based’ approach to communities, where interactions are merely treated as fixed covariates affecting distribution.

As a more explicit description of interactions among species, the network approach offers a convenient representation of communities. Species are represented as nodes and interactions by links. To date, studies of network diversity have mostly been concerned with the distribution of interactions within locations, and less so with the variation among locations (????). There is however ample evidence that interaction networks vary in space and time (?????). Metacommunity theory provides explanation for the variation in the distribution of the different nodes (??), but there is no explanation to the joint variation of node and link occurrences. We urgently need a conceptual framework to formalize these observations.

Given the historically different approaches to modelling the distributions of species vs. interactions, there is an evident need to bring the two together. Here, we offer an integrated approach, adopting the view that community structure is best represented as

90 an ecological network of interactions. Based on this idea, we propose a new description
 91 of the basic concept of the ecological niche, now integrating the effect of the environment
 92 on species distribution and on the ecological interactions among them. Starting from this
 93 redefined concept, we develop a quantitative theory to explain turnover of interactions
 94 in space and time. We first present the conceptual framework, and then formalize it
 95 mathematically, using a probabilistic model to represent the sampling of the regional pool
 96 of interactions. At the level of species pairs, the statistical approach could be conceived
 97 as an interaction distribution model. At the community level, the approach provides a
 98 likelihood-based method to compare different hypotheses of network turnover. We apply
 99 this novel framework to a large data set on host–parasite interactions across Europe and
 100 find that variation of the environment causes turnover of both species and interactions.
 101 The network structure changes systematically across the latitudinal gradient, with a peak
 102 of connectance at intermediate latitude.

103 **The two dimensions of community structure**

104 The problem of community assembly is often formulated as *how are species sampled from*
 105 *a regional pool to constitute a local community (?)*? This question could be rewritten
 106 to address the problem of network assembly, as *how do samples from a regional pool of*
 107 *interactions constitute a local interaction network*? An illustration of this problem for a
 108 food web is provided in Figure 1. The metaweb represents potential interactions among
 109 all species that could be found in a given area. In this particular case, there are 275
 110 nodes, and 1173 links among plants (52 nodes), herbivores (96 nodes), and parasitoids
 111 (127 nodes) from Northern Europe. An instance of a local community is also illustrated,
 112 with 45 nodes and 93 interactions. Only 55.0% of all potential interactions are realized,
 113 revealing the stochastic nature of ecological interactions. Our objective here is to provide

114 a conceptual framework to explain the sampling of the regional pool of interactions, along
115 with a quantitative method to predict it. The problem could be formalized sequentially
116 by first understanding why only a fraction of the species are co-occurring locally and
117 second why these species do or do not interact.

118 There are multiple causes of spatial turnover in community composition. The first
119 and most-studied driver is the effect of variation in the abiotic environment on species
120 performance. Combined with specific responses in demography, it generates variation
121 among sites by selecting the locally fittest species (?). Stochasticity plays an additional
122 role, either colonization and extinction events (?) are inherently unpredictable, or because
123 of strong non-linear feedbacks in community dynamics generating alternative transients
124 and equilibria (??). Analyses of community turnover are usually performed with data
125 represented in a table with rows corresponding to sites (or measurements) and columns
126 to species. Metrics of beta diversity quantify the variance of this community data (?).
127 Traditional approaches rely on measures of dissimilarity among communities, such as the
128 Jaccard or Bray–Curtis indices. A more recent approach decomposes the total variation of
129 the community data into species and site contributions to beta diversity (?), and further
130 partition into dissimilarity due to changes in species richness and those due to actual
131 species turnover (?). Even though these methods compare whole lists of species among
132 sites or measurements, they remain fundamentally 'species- based', since they report the
133 variation within columns. None of them explicitly considers the variation of associations
134 (i.e., of pairs or higher- order motifs – ?).

135 The "niche" is by far the dominant concept to explain species distributions and com-
136 munity assembly, from the local to the global scale. Following ?, the niche is viewed as
137 the set of environmental conditions allowing a population to establish and persist (see
138 also ?). Community turnover arises as a result of successive replacement of species along
139 an environmental gradient, in agreement with the Gleasonian view of communities (?).

140 The concept is straightforward to operationalize with species distribution models, as it
141 maps naturally on available distributional and environmental data. In consequence, a
142 vast array of statistical tools have been developed to implement it (e.g. BIOMOD ?,
143 MaxEnt ?). It is however much harder to account for ecological interactions with this
144 approach (?). As such, these interactions are often viewed as externalities constraining
145 or expanding the range of environmental conditions required for a species to maintain a
146 viable population (???)

147 As mentioned above, in the network approach to community structure, species and
148 interactions are represented by nodes and links, respectively. Associations can also be
149 represented by matrices in which entries represent the occurrence or intensity of interac-
150 tions among species (rows and columns). Network complexity is then computed as the
151 number of interactions (in the case of binary networks) or interaction diversity (in the
152 case of quantitative networks, ?). Variability in community structure consequently arises
153 from the turnover of species composition, along with turnover of interactions among pairs
154 of species. The occurrence and intensity of interactions could vary because of the envi-
155 ronment, species abundance, and higher-order interactions (?). Variation in community
156 composition was found independent of the variation of ecological interactions, suggesting
157 that these two components of network variability respond to different drivers (?).

158 Interestingly, the ecological network literature also has its own 'niche model' to po-
159 sition a species in a community (?). The niche of a species in this context represents
160 the multidimensional space of all of its interactions. Each species is characterized by a
161 niche position, an optimum and a range over 3 to 5 different niche axes (??). The niche
162 model of food web structure and its variants have successfully explained the complexity
163 of a variety of networks, from food webs to plant-pollinator systems (???). This con-
164 ceptual framework is, however, limited to local communities, and does not provide any
165 explanation for the turnover of network structure along environmental gradients.

The integrated niche

A more integrative description of the niche is key to understand spatial and temporal turnover in community structure. Despite several attempts to update the concept of the ecological niche, ecologists have not moved far past the “n-dimensional hypervolume” defined by Hutchinson. Despite its intuitive interpretation and easy translation into species distribution models (??), the concept has been frequently criticized (???), and several attempts have been made to expand and improve it (?????).

Part of the problem surrounding the niche concept has been clarified with the distinction between Eltonian and Grinnellian definitions (?). The Grinnellian dimension of the niche is the set of environmental conditions required for a species to maintain a population in a location. The Grinnellian niche is intuitive to apply, and constitutes the conceptual backbone of species distribution models. The Eltonian niche, on the other hand, is the effect of a species on its environment. This aspect of the niche is well known by community ecologists, but is trickier to turn into predictive models. Nonetheless, the development of the niche model of food web structure (?) and its parameterization (?) made it more operational.

These perspectives are rather orthogonal to each other, which has resulted in considerable confusion in the literature (?). ? attempted to reconcile with the following definition: “[The niche is] the joint description of the environmental conditions that allow a species to satisfy its minimum requirements so that the birth rate of a local population is equal to or greater than its death rate along with the set of per capita effects of that species on these environmental conditions”. Their representation merges zero-net growth isoclines delimiting the Grinnellian niche (when does the population persists) with impact vectors delimiting the Eltonian niche (what is the per-capita impact). While this representation has been very influential in local-scale community ecology (the resource-ratio theory of coexistence, ?), it remains impractical at larger spatial scales because of the difficulties

in measuring it. The absence of any mathematical representation of the niche that could easily be fitted to ecological data may explain why biogeographers are still struggling to develop species distribution models that also consider ecological interactions.

We propose to integrate the two perspectives of the niche with a visual representation of both components (Fig. 2). The underlying rationale is that, in addition to the environmental constraints on demographic performance, any organism requires resources to sustain its metabolic demands and reproduction. Abiotic environmental axes are any non-consumable factors affecting the demographic performance of an organism. Alternatively, the resource axes are traits of the resources that allow interactions with the consumer. The niche should therefore be viewed as the set of abiotic environmental factors (the Grinnellian component) along with the set of traits (the Eltonian component) that allow a population to establish and to persist at a location. Accordingly, each species can be characterized by an optimal position in both the environmental (x-axis) and the trait (y-axis) plane. The integrated niche is then the hypervolume where interactions can occur and sustain a population. This approach radically changes the representation of the niche, putting the distributions and ecological interactions of species in the same formalism.

The limits of the niche axes could be independent of each other (as in the example in Fig. 2), or they could interact. For instance, the optimal prey size for predatory fishes could decline with increasing temperature (?), which would make diet boundaries functions of the environment. Alternatively, we could also consider that the growth rate of the predator changes with the size of its prey items, thereby altering the environmental boundaries.

A probabilistic representation of interaction networks in space

We now formalize the integrated niche with a probabilistic approach to interactions and distributions. We seek to represent the probability that an interaction between species i and j occurs at location y . We define L_{ijy} as a stochastic variable, and are looking at the probability that this event occurs, $P(L_{ijy})$. The occurrence of an interaction is dependent on the co-occurrence of species i and j . This argument might seem trivial at first, but the explicit consideration of this condition in the probabilistic representation of ecological interactions will prove fundamental to understand their variation. We define X_{iy} as a stochastic variable representing the occurrence of species i at location y . The quantity we seek to understand is the probability of a joint event, conditional on the set of environmental conditions E_y :

$$P(X_{iy}, X_{jy}, L_{ijy} | E_y) \quad (1)$$

Or simply said, the probability of observing both species i and j plus an interaction between i and j given the conditions E_y . This probability could be decomposed into two parts using the product rule of probabilities:

$$P(X_{iy}, X_{jy}, L_{ijy} | E_y) = P(X_{iy}, X_{jy} | E_y) P(L_{ijy} | \mathbf{T}, X_{iy}, X_{jy}, E_y) \quad (2)$$

The first term on the right-hand side of the equation is the probability of observing the two species co-occurring at location y . It corresponds to the Grinnellian dimension of the niche. The second term represents the probability that an interaction occurs between species i and j , given their set of traits T and that they are co-occurring. Above, we refer to this entity as the "metaweb" and it corresponds to the Eltonian dimension of the

niche. Below, we will see how this formalism can be directly fitted to empirical data. But before turning to an application, we will discuss the interpretation of different variants of these two terms.

Variants of co-occurrence

There are several variants to the co-occurrence probability, representing different hypotheses about temporal and spatial variation in network structure (see the explicit formulations in Table 1). The simplest model relates the probability of co-occurrence directly to the environment, $P(X_{iy}, X_{jy}|E_y)$. In this situation there are no underlying assumptions about the ecological processes responsible for co-occurrence. It could arise because interactions do constraint distribution (??) or, alternatively, because of environmental requirements shared between i and j . In the former case, species are not independent of each other and the conditional occurrence must be accounted for explicitly, $P(X_{iy}, X_{jy}|E_y) = P(X_{iy}|E_y, X_{jy})P(X_{jy}|E_y)$. In the latter case, species are independent and only the marginal occurrence must be accounted for, $P(|E_y) = P(X_{iy}|E_y)P(X_{jy}|E_y)$.

The co-occurrence probability itself could depend on ecological interactions. This should be viewed as the realized component of the niche (i.e. the distribution when accounting for species interactions). Direct pairwise interactions such as competition, facilitation, and predation have long been studied for their impact on co-distribution (e.g. ???). Second- and higher-order interactions (e.g. trophic cascades) could also affect co-occurrence. Co-occurrence of multiple species embedded in ecological networks is a topic of its own, however, and is influenced by both network topology and species richness (?). Not only direct interactions do influence co-occurrence, but indirect interactions do as well (e.g. plant species sharing an herbivore could repulse each other in space). The impact of direct interactions and first-order indirect interactions on co-occurrence tends to vanish with increasing species richness in the community. Further, co-occurrence is

also influenced by the covariance of interacting species to an environmental gradient
 (?). Because of the complexity of relating co-occurrence to the structure of interaction
 networks, we will here focus on the variation of interactions and not on their distribution,
 and leave this specific issue for the Perspectives and future research.

Variants of the metaweb

There are also variants of the metaweb. First, most documented metawebs have thus far
 considered ecological interactions to be deterministic, not probabilistic (e.g. ??). Species
 are assumed to interact whenever they are found together in a location, independent of
 their local abundance and the local environment. In other words, $P(L_{ijy}|X_{ijy} = 1) = 1$
 and $P(L_{ijy}|X_{ijy} = 0) = 0$. This approach might be a reasonable approximation if the
 scale of sampling and inference is so large that the probability of observing at least one
 interaction converges to unity. In this scenario, network variation solely arises from species
 distributions.

Ecological interactions could also vary with the environment, so that $P(L_{ijy}|E_y)$.
 Although it is not common to see a conditional representation of pairwise ecological
 interactions, experimental studies have revealed them to frequently be sensitive to the
 environment. For instance, (?) showed that predation risks of shorebirds vary at the
 continental scale, decreasing from the south to the north. It is also common to see
 increasing top-down control with temperature (e.g. ??). Effects of the environment on
 interactions also propagate up the community and influence network structure (??).

Application: continental-scale variation of host-parasite community structure

We now turn to an illustration of our framework with the analysis of an empirical dataset of host–parasite networks sampled throughout the south–north environmental gradient in continental Europe. The focal system consists of local food webs of willows (genus *Salix*), their galling insects, and the parasitoids of these gallers. Targeting this system, we ask: i) how much does network structure vary across the gradient, and ii) what is the primary driver of network turnover across the gradient?

Data

Communities of willows, gallers, and parasitoids are species-rich and widely distributed, with pronounced variation in community composition across space. The genus *Salix* includes over 400 species, most of which are shrubs or small trees (?). The genus is common in most habitats across the Northern Hemisphere (?). Willows support a highly diverse community of herbivorous insects, with one of the main herbivore groups being gall-inducing sawflies (Hymenoptera: Tenthredinidae: Nematinae: Euurina (?)). Gall formation is induced by sawfly females during oviposition, and includes marked manipulation of host-plant chemistry by the galler (?). The parasitoid community includes nearly 100 species belonging to 17 insect families of four orders (?). These encompass two main types: inquiline larvae (Coleoptera, Lepidoptera, Diptera, and Hymenoptera) feed primarily on gall tissue, but typically kill the galler larva in the process, while parasitoid larvae (representing many families in Hymenoptera) kill the galler larvae by direct feeding (?). In terms of associations between the trophic levels, phylogeny-based comparative studies have demonstrated that galls represent “extended phenotypes” of the gallers, meaning that gall form, location, and chemistry is determined mainly by the

galling insects and not by their host plants (?). Because galler parasitoids have to penetrate a protective wall of modified plant tissue in order to gain access to their victims, gall morphology has been inferred to strongly affect the associations between parasitoids and hosts (?). Thus, the set of parasitoids attacking each host is presumably constrained by the form, size, and thickness of its gall.

Local realizations of the willow–galler–parasitoid network were reconstructed from community samples collected between 1982 and 2010. During this period, willow galls were collected at 370 sites across Central and Northern Europe. Sampling was conducted in the summer months of June and/or July, i.e., during the later stages of larval development. Galler species were identified on the basis of willow host species and gall morphology, as these are distinct for each sawfly species. At each site, galls were randomly collected from numerous willow individuals in an area of about 0.1–0.3 km^2 . Some sites were visited more than once, with a total of 641 site visits across the 370 sites. GPS coordinates were recorded for each location; for our analyses, the present day annual mean temperature and precipitation were obtained from WorldClim using the R package raster (?). While other covariates could have also been considered, these variables are likely representative of the most important axes of the European climate, and more easily interpretable than reduced variables obtained by e.g. principal component analysis.

The methods used for rearing parasitoids from the galls have been previously described by ?. In brief, galls were opened to score the presence of galler or parasitoid/inquiline larvae. Parasitoid larvae were classified to preliminary morphospecies, and the identity of each morphospecies was determined by connecting them to adults emerging after hibernation. The galls were reared by storing single galls in small glass tubes (?). Hibernation of galls containing parasitoids took place either within the glass tubes or between blotting paper in flowerpots filled with clay granulate or a mixture of peat dust and sand. These pots were stored over the winter in a roof garden and/or in a climatic chamber. In most

cases, the matching of larval morphospecies with adult individuals emerging from the rearings allowed the identification of the parasitoids to the species level. Nonetheless, in some cases, individuals could only be identified to one of the (super)families Braconidae, Ichneumonidae, and Chalcidoidea. This was particularly the case when only remains of faeces, vacant cocoons of parasitoids, and/or dead host larvae were found, as was the case when parasitoids had already emerged from the gall. As a result, the largest taxon in the data set, “Chalcidoidea indeterminate”, represents a superfamily of very small parasitoids that are hard to distinguish.

In total, 146,622 galls from 52 *Salix* taxa were collected for dissection and rearing. These galls represented 96 galler species, and yielded 42,133 individually-identified parasitoids. Of these, 25,170 (60%) could be identified to the species level. Overall, 127 parasitoid taxa were distinguished in the material. Data on host associations within subsets of this material have been previously reported by (?) and (?). The current study represents the first analysis of the full data set from a spatial perspective.

Analysis

Computing the probability of observing an interaction involves fitting a set of binomial models and collecting their estimated probabilities. For the sake of illustration, we considered second-order generalized linear models – whereas more flexible fitting algorithms (e.g. GAM or Random Forest) could equally well be used, as long as the algorithm can estimate the probability for each observation. The data consist of a simple (albeit large and full of zeros) table with the observation of each species, X_{iy} and X_{jy} , their co-occurrence, X_{ijy} , the observation of an interaction L_{ijy} , and environmental co-variables E_y . Thus, there is one row per pair of species per site. We considered that an absence of a record of an interaction between co-occurring species at a site means a true absence (see below for a discussion on this issue).

355 We compared three models for the co-occurrence probability. The first one directly
 356 models the co-occurrence probability conditional on the local environment, $P(X_{iy}, X_{jy}|E_y)$
 357 (models are listed at Table 1 and 2). Hence, this model makes no assumptions about the
 358 mechanisms driving co- occurrence for any given environment, and instead uses the in-
 359 formation directly available in the data. It thereby indirectly accounts for the effect of
 360 interactions on co-occurrence, if there is any. The second model considers independent
 361 occurrence of species. In this case, we independently fit $P(X_{iy}|E_y)$ and $P(X_{jy}|E_y)$, then
 362 we take their product to derive the probability of co-occurrence. This model should be
 363 viewed as a null hypothesis with respect to the first model, since a comparison between
 364 the respective models will reveal if there is significant spatial association of the two species
 365 beyond a joint response to the shared environment (?). Finally, the third model assumes
 366 that the probability of co-occurrence is independent of the environment and thus constant
 367 throughout the landscape. In other words, $P(X_{iy}, X_{jy})$ is obtained by simply counting the
 368 number of observed co-occurrences divided by the total number of observations. Thus,
 369 the comparison between the first and third model allows us to test the hypothesis that
 370 co-occurrence is conditional on the environment. Whenever the environment was included
 371 as a covariate in the GLM, we considered a second- order polynomial response for both
 372 temperature and precipitation. There are consequently 5 parameters for the first model
 373 when fitting a given pair of species, 10 parameters for the second, and only 1 for the third
 374 model.

375 Following the same logic, we compared three models of the interaction probability. The
 376 first model conditions the interaction probability on the local environmental variables,
 377 $P(L_{ijy}|X_{iy}, X_{jy}, E_y)$. Consequently, the model was fit to a subset of the data where the
 378 two species co-occur. The second model fits the interaction probability independently of
 379 the local environmental variables, $P(L_{ijy}|X_{iy}, X_{jy})$. It corresponds to the number of times
 380 the two species were observed to interact when co-occurring, divided by the number of

times that they co-occurred. The third model is an extreme case performed only to test the hypothesis that if two species are found to interact at least once, then they should interact whenever they co-occur, $P(L_{ijy}|X_{iy}, X_{jy}) = 1$. While not necessarily realistic, this model tests an assumption commonly invoked in the representation of local networks from the knowledge of a deterministic metaweb. There are consequently 5 parameters for the first model, a single parameter for the second model and no parameter to evaluate for the third model (where the interaction probability is fixed by the hypothesis).

The different models were fitted to each pair of species and the predicted probabilities were recorded. The joint probability $P(L_{ijy}, X_{iy}, X_{jy})$ was then computed from Eq. 2, and the likelihood of each observation was computed as $\text{mathcal{L}}(\theta_{ijy}|D_{ijy}) = P(L_{ij}, X_{iy}, X_{jy})$ if an interaction was observed, and as $\text{mathcal{L}}(\theta_{ijy}|D_{ijy}) = 1 - P(L_{ijy}, X_{iy}, X_{jy})$ if no interaction was observed. The log-likelihood was summed over the entire dataset to compare the different models by AIC. Not surprisingly, there was a very large number of species pairs for which this model could not be computed, as they simply never co-occurred. For these pairs, we have no information of the interaction probability, and they were consequently removed from the analysis. The log-likelihood reported across the entire dataset was summed over all pairs of species observed to co-occur at least once. Interactions between the first (Salix) and second (gallers) trophic layers and those between the second and third (parasitoids) were considered separately. Finally, we used the full model (in which both co-occurrence and the interaction are conditional on the environment) to interpolate species distributions and interaction probabilities across the entire European continent. We reconstructed the expected network for each location in a 1km X 1km grid and computed the probabilistic connectance following (?).

All of the data are openly available in the database mangal (?) and all R scripts for querying and pre-processing the data, along with the analysis, are provided in the Supplementary material.

Results

Despite the extensive sampling, many pairs of species were observed to co-occur only a few times. This made it hard to evaluate interaction probabilities with any reasonable confidence interval. Thus, we start with an example of a single pair of species selected because of its high number of co-occurrences ($N_{ij} = 38$): *Phyllocolpa prussica* and *Chrysocharis elongata*. These two fairly abundant species, were observed $N_i = 49$ and $N_j = 121$ times, respectively, across the 370 sites, and they were found to interact with a marginal probability $P(L_{ij}) = 0.55$, which means they interacted at 21 different locations. Here, a comparison of model fit (Table 1) reveals that the interaction probability conditional on co-occurrence does not better explain their distribution (Model 1 vs Model 2). Moreover, when the two species co-occur, the occurrence of the interaction was insensitive to the environment (Model 2 vs Model 3). Alternatively, climatic variables significantly impacted co-occurrence (Model 3 vs Model 4). The neutral model performed worse than the non-random co-occurrence model (Model 3 vs Model 6). The full model revealed that the greatest interaction probability occurred at intermediate temperature and precipitation, simply because this is where the two species most frequently co-occur (Fig. 3). The probabilities of co-occurrence and interaction can be represented in space, where we found that the highest interaction probability occurred in central Europe (Fig. 4).

We evaluated each model for all pairs of species in order to better understand the large-scale drivers of network turnover. The results were highly consistent among trophic layers (Salix-gallers and gallers-parasitoids; Table 2). Across all pairs of species, the conditional representation of interactions performed better than the marginal one (Model 1 vs Model 2); that is, interactions did not occur systematically whenever the two species were found co-occurring. Hence, in addition to species turnover, the stochastic nature of interactions contributes to network variability. In total, we recorded 1,173 pairs of

interactions, only 290 of which occurred more than 5 times. Out of these 290 interactions, 143 were systematically detected whenever the two species co-occurred. In the instances when species co-occurred, the two environmental variables considered proved relatively poor predictors of their interactions (Model 2 vs Model 3). Not surprisingly, for both types of interactions (Salix–galler and galler–parasitoid), the likelihood increased when the environment was considered. However, the extra number of parameters exceeded the gain in likelihood and inflates AIC. Therefore, the most parsimonious model excluded the effect of the environment. On the basis of log-likelihoods only, co-occurrence was non-neutral for both Salix–galler interactions and galler–parasitoid interactions. Thus, according to AIC, the best model was the one of non-random co-occurrence (Model 3 vs Model 6) for both types of interactions.

The approach we present not only has implications for understanding the biogeography of pairwise interactions and interaction networks, but also the quality of the evaluation of metawebs. We investigated the reliability of the estimated metaweb across the entire dataset with summary statistics of species co-occurrence. As mentioned above, across the 17,184 potential pairs of species, only 1,173 pairs interacted in at least a single location, yielding a connectance of 0.068. However, only 4,459 pairs of species were found co-occurring at least once across all locations. There are consequently 12,725 gaps of information in the metaweb (74.1% - see Fig. 5). As we cannot know whether the non-co-occurring species would indeed interact when found together, a more appropriate estimate of connectance would be $C = 1173/4459 = 0.263$. This result reveals that the evaluation of the sampling quality of ecological networks is a problem on its own and well worth further attention.

Once we had selected the best model (Model 3, Table 2), we used it to reconstruct the expected species richness, along with the most likely network for each location. By this approach, we could map the expected distribution of network properties across Europe

(Fig. 6). For simplicity, we chose to consider connectance as descriptor of network configuration, as this metric can be easily computed from probabilistic networks (?) and is also a good proxy for many other network properties (?). Overall, we find a peak in *Salix*, gallers and parasitoid diversity in northern Europe. The expected number of interactions roughly follows the distribution of species richness, but accumulates at a rate different from species numbers. Connectance also peaks in northern Europe (Fig. 6).

Interpretation

We have proposed that the representation of community structure and its variation in space and time is best represented with the formalism of ecological networks because both the distribution of species and their interspecific interactions can then be accounted for. We consequently revised the niche concept in order to integrate both the abiotic and the biotic components of the niche that are susceptible to vary over time and space. This integrated niche was represented visually with an ordination of species into an environmental space and a trait space. The fundamental niche of a species is represented as the set of environmental conditions and resources that allow a species to be maintained in a location, thereby integrating the Eltonian and the Grinnellian components of the niche. We then translated the concept mathematically by investigating the probability of the joint occurrences of species and their interaction, which should be interpreted as an interaction distribution model. We used this approach to characterize the turnover of the structure of ecological interactions in a tri-trophic network across Western Europe and found that the primary driver of variation is the turnover in species composition. To our knowledge, this is the first continental-wide analysis of the drivers of network structure from empirical data (see ??).

Applying the framework to our large data set on host–parasite interactions across Eu-

rope revealed key features in the interaction between *Salix* taxa, their herbivores, and the natural enemies of these herbivores. Consistent with a general increase in the diversity of *Salix* towards boreal areas (?), overall species richness of the networks increased towards the north. The distribution of *Salix* species richness largely matched those of gallers and parasitoids. These observations within Europe are also matched by the ones found at a global scale for *Salix* (???) and sawflies (?). Species richness was originally presumed to show a similar "reversed latitudinal gradient" for a common group of parasitic wasps, the Ichneumonidae, but this observation has been challenged by findings of rather high ichneumonid diversity in the tropics (?). However, the ichneumonid subfamilies specifically associated with sawflies (Ctenopelmatinae, Tryphoninae) are clearly less diverse in the south.

Exactly what processes are responsible for the distribution of species richness at different trophic levels is yet to be established (but see e.g. ???), but as a net outcome of different latitudinal trends across trophic levels, the distribution of co-occurrence and therefore of potential interactions differed between the first and second link layers. The correlation between the expected *Salix* and gallers richness was 0.73, while it was 0.58 between gallers and their parasitoids. The ratio of herbivore to *Salix* species is essentially constant across Europe, whereas each herbivore species is potentially attacked by a richer enemy community higher latitudes (i.e. faces higher vulnerability). Consequently, overall connectance peaks in Northern Europe (Fig. 6).

In terms of species interacting with each other, our analysis suggests that the environment leaves a detectable imprint on species co-occurrence, but only a slight mark on the occurrence of realized links among species in a specific place: the probability of finding a given combination of species at a higher and a lower trophic level at the same site was clearly affected by the environment, whereas the probability of observing an interaction between the two was not detectably so. This applies to the example species *Phyllocolpa*

509 *prussica* and *Chrysocharis elongata* (Figs 2-3), but also to all species pairs more generally.
510 For the example species pair, the full model revealed that the interaction probability is
511 highest at intermediate temperature and precipitation, simply because this is where the
512 two species co-occur most often. This does not imply that species will always interact
513 when they meet – although this is a basic assumption in most documented metawebs to
514 date (e.g. ??). Rather, an interaction is a stochastic process whose probability is also
515 influenced by the probability with which species co-occur. What we cannot reliably know
516 is how this stochasticity splits into two sampling processes – i.e., the extent to which a
517 species at the higher trophic level runs into a species at the lower level co-occurring at
518 the site, and the extent to which this interaction is detected by an observer collecting a
519 finite sample. Future work will be required to document the relative importance of these
520 two sources of uncertainty in the occurrence of interactions.

521 Perspectives

522 Evidence that the structure of ecological networks does vary across habitats (e.g. ??),
523 over environmental gradients ? and in time (?) is accumulating rapidly. It is not
524 clear however to what extent the turnover of network structure is driven by a systematic
525 change in species composition or of pairwise interactions (??). The model comparison of
526 the host-parasite interactions revealed that most of the turnover is driven by a species-
527 specific response to the environment, impacting species richness, and that co-occurrence
528 was mostly neutral. Further, the occurrence of interactions in presence of the host and
529 parasite is highly stochastic, but not predictable according to the variables we considered.
530 We know that interactions vary with the environment in other systems, for instance,
531 herbivory (?), predation (??) are often found increasing with temperature, resulting in
532 spatial variation of trophic cascades (?). What remains unclear, however, is the extent

533 to which such variation is driven by a turnover of species composition along gradients,
534 or a turnover of the interactions. Here we found that interactions vary substantially, but
535 not predictably along the annual temperature and the precipitation gradient. Perhaps
536 we have not found a strong signal of the effect of the environment on the occurrence
537 of interactions because we had wrong covariates. It was indeed found previously for a
538 similar system that habitat characteristics are the primary drivers of interactions (?). New
539 investigations with other systems will be required to challenge this result. Documenting
540 the relationship between the occurrence of interactions at the continental scale and the
541 environment is critical to understand how trophic regulation at large spatial scales do
542 influence community dynamics and ecosystem functioning (?).

543 We restricted our framework to the effect of co-occurrence on ecological interactions
544 and neglected the inverse of the problem. We did not investigated in depth the drivers
545 of co-occurrence and simply took it for granted from the data. Co-occurrence was indeed
546 many times significantly different from the expectation of independent species distri-
547 bution. It thus raises the question that, once accounting for the species-specific effect
548 of the environment on distribution, are there significant effects of interactions on co-
549 occurrence? We could rephrase this problem asking if the fundamental niche differs from
550 the realized niche, and how it applies to our framework. For example, we have considered
551 above simply the co-occurrence probability, $P(X_{iy}, X_{jy}|E_y)$, which could be expanded
552 as $P(X_{iy}|X_{jy}, E_y)P(X_{jy}|E_y)$. The marginal occurrence probability, $P(X_{jy}|E_y)$, could be
553 considered as a species distribution model taking into account the interaction between
554 these species after some re-arrangement of Eq. 2. This derivation would however criti-
555 cally depend on a strong a priori expectation of the conditional probability of observing
556 a species given the distribution of the other species. This assumption seems reasonable
557 for some situations, such as a parasitoid species that requires a host to develop. On the
558 other hand, we found the strength of this association is often rather weak if not neutral

(with the example pair for instance analysed at Table 1). The lack of an association could simply arise when the parasitoid is generalist enough so that it is not constrained to track the distribution of its host (?). At present, there is only indirect support to the hypothesis that interacting species are conditionally distributed and it should be the topic of more specific hypothesis testing. The impact of ecological interactions on the distribution of co-occurrence has been the topic of many publications since ? seminal study on competition and the checkerboard distribution, but only recently pairwise approaches received attention (?). It therefore remains unclear if two interacting species are more closely associated in space because most approaches based on null models consider community-level metrics (e.g. ?), such as the C-score, thereby making it hard to evaluate if specific interactions do indeed affect co-occurrence. The expansion of the framework we described to account for the difference between the realized and the fundamental niche will therefore require further investigation of the impact of interactions on co-occurrence.

Ecological networks are known to be extremely sparse, *i.e.* having far more absences of interactions than they have interactions. These absences of interactions, however, can come from different sources. The fact that unequal sampling at the local scale can affect our understanding of network structure is well documented (?). However, in a spatial context, some interactions may be undocumented because the species involved have never been observed to co-occur. Although these are reported as a lack of interactions, in actuality we can not make inference about them since they have never been observed: it remains possible that this interaction may happen should the two species co-occur. A second category of absences of interactions are those that are reported after multiple observations of species co-occurrence. However, so as to have a confidence in the fact that the probability of an interaction is low, extensive sampling (that is, several records of co-occurrence) is needed. Generally, our confidence that the interaction is indeed impossible will increase with the number of observations of the species pair. Seeing that this is

585 essentially a Bernoulli process (the probability that the species will interact given their
586 presence), the breadth of the confidence interval is expected to saturate after a fixed
587 number of observations, which can be set as a threshold above which a species pair has
588 finally been observed "often enough".

589 Conclusion

590 Understanding the drivers of the spatial variation in network structure is a key problem
591 to solve in order to anticipate the impacts of global changes on biodiversity. Our rep-
592 resentation of the spatial variation of community structure presents a new approach for
593 the study of the biogeography of ecological networks. We see the following key challenges
594 and opportunities ahead in this exciting area of research:

595 **1. New generation of network data.** The investigation of the spatial variation
596 of network structure will require high quality and highly replicated network data. We
597 have investigated one the most comprehensive spatial network datasets we are aware of
598 and nonetheless found immense gaps of knowledge in its resolution. Species richness
599 accumulates much faster than observations of ecological interactions (?). Each pair of
600 species must be observed several times to have reliable estimates of their interaction
601 probability.

602 **2. Estimation of the reliability of interactions.** We need quantitative tools to
603 estimate the confidence interval around an estimate of interaction probability, as well as
604 some estimation of the rate of false absences. Bayesian methods are promising to that
605 end because we could use information on the target species (e.g. if they are known as
606 specialists or generalists) to provide prior estimates of the interaction probability.

607 **3. From interaction probabilities to a distribution of network properties.**
608 Metrics are available to analyse the structure of probabilistic networks (?). These met-

rics are useful as first approximation, but they assume independence among interactions. This might not be the case in nature because of the role of co-occurrence and shared environmental requirements. We also need to better understand the distribution of network properties arising from probabilistic interactions.

4. Investigation of the environmental-dependence of ecological interactions.

There is evidence that interactions can vary in space, but this problem has not been investigated in a systematic fashion. The paucity of the data currently available precludes an extensive analysis of this question.

5. Effects of ecological interactions on co-occurrence. We have intentionally omitted the feedback of ecological interactions on co-occurrence in this framework. As abundance can impact the occurrence of interactions, and inversely since interactions do impact abundance (?), we could reasonably expect the same for co-occurrence. Theory does exist for simple three-species modules (?), but the extension to entire co- occurrence networks will prove critical, especially given the interest in using co-occurrence to infer ecological interactions (??).

Acknowledgements

This is a contribution to the working groups *Networks over ecological gradients* (Santa Fe Institute) and *Continental-scale variation of ecological networks* (Canadian Institute for Ecology and Evolution).

Table 1: Summary of model comparison for the interaction between the leaf galler *Phyllocolpa prussica*) and the parasitoid *Chrysocharis elongata*

#	Metaweb model	Co-occurrence model	LL	npars	AIC
1	$P(L_{ijy})$	$P(X_{iy}, X_{jy} E_y)$	-65.5	6	143
2	$P(L_{ijy} X_{iy}, X_{jy})$	$P(X_{iy}, X_{jy} E_y)$	-65.7	6	143.4
3	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy}, X_{jy} E_y)$	-65.6	10	151.3
4	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy}, X_{jy})$	-84.5	6	183
5	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy})P(X_{jy})$	-80.7	7	173.4
6	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy} E_y)P(X_{jy} E_y)$	-68.8	15	167.6

Table 2: Summary of model comparison for the interaction across all pairs of salix, gallers and parasitoids.

Interaction	#	Metaweb model	Co-occurrence model	LL	npars	AIC
Plant-Galler	1	$P(L_{ijy})$	$P(X_{iy}, X_{jy} E_y)$	-6137.8	7170	26615.6
	2	$P(L_{ijy} X_{iy}, X_{jy})$	$P(X_{iy}, X_{jy} E_y)$	-5947.2	7170	26234.3
	3	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy}, X_{jy} E_y)$	-5939.8	11950	35779.6
	4	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy}, X_{jy})$	-7871.9	8365	32473.8
	5	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy})(X_{jy})$	-6639.9	7170	27619.9
	6	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy} E_y)P(X_{jy} E_y)$	-7123.2	17925	50096.4
Galler-Parasitoid	1	$P(L_{ijy})$	$P(X_{iy}, X_{jy} E_y)$	-21397.6	18846	81963.1
	2	$P(L_{ijy} X_{iy}, X_{jy})$	$P(X_{iy}, X_{jy} E_y)$	-21105.2	18846	81378.5
	3	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy}, X_{jy} E_y)$	-20881.1	31410	107042.1
	4	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy}, X_{jy})$	-23728.3	21987	93152.6
	5	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy})P(X_{jy})$	-23509.4	18846	86186.8
	6	$P(L_{ijy} X_{iy}, X_{jy}, E_y)$	$P(X_{iy} E_y)P(X_{jy} E_y)$	-20990	47115	139900.1

Figure legends

Figure 1

Non-random sampling of the metaweb. Network assembly can be viewed as a sampling process of the regional pool of potential interactions. Species (indicated by colored nodes) are sampled first, and among the species found in the local network, only some interactions (indicated by colored links) occur. We characterize these sampling processes with the quantitative framework proposed in this paper. As a concrete illustration of metaweb sampling, we here show a local interaction network among *Salix*, gallers, and parasitoids. The metaweb was constructed by aggregating observed interactions across 370 local networks.

Figure 2

Visual representation of the integrated niche. In biogeography, the niche is considered the set of environmental conditions where the intrinsic growth rate r is positive (?). The horizontal axis represents an environmental gradient impacting the growth of the focal species (in red). The location of each species along this gradient represents their optimum, and the vertical dotted lines represent the limits of the Grinnellian niche of the focal species. In food web ecology, the Eltonian niche represents the location of a species in the food web, as determined by its niche position (n) and its niche optimum (c). The vertical axis represents a niche gradient, presumably a trait such as body size. The location of each species along this gradient represents their niche position. The focal species will feed on prey species niche locations within a given interval around the optimum, represented by the horizontal lines. The integrated Grinnellian and Eltonian niche corresponds to the square in the middle where an interaction is possible owing to a match of traits and spatial distribution. According to our probabilistic framework, the

central square represents the area where the joint probability of observing interactions and co-occurrence is positive.

Figure 3

Probabilistic representation of the interaction probability between a leaf galler (*Phyllocolpa prussica*) and a parasitoid (*Chrysocharis elongata*) across an annual average temperature and an annual precipitation gradient. The representation is based on predictions from model 3 (see Table 1). In the left panel, open circles represent the absence of both species, whereas closed circles represent co-occurrence and plus signs the occurrence of only one of the two species. In the other two panels, open circles represent co-occurrence but an absence of interaction and the closed circles represent the occurrence of an interaction.

Figure 4

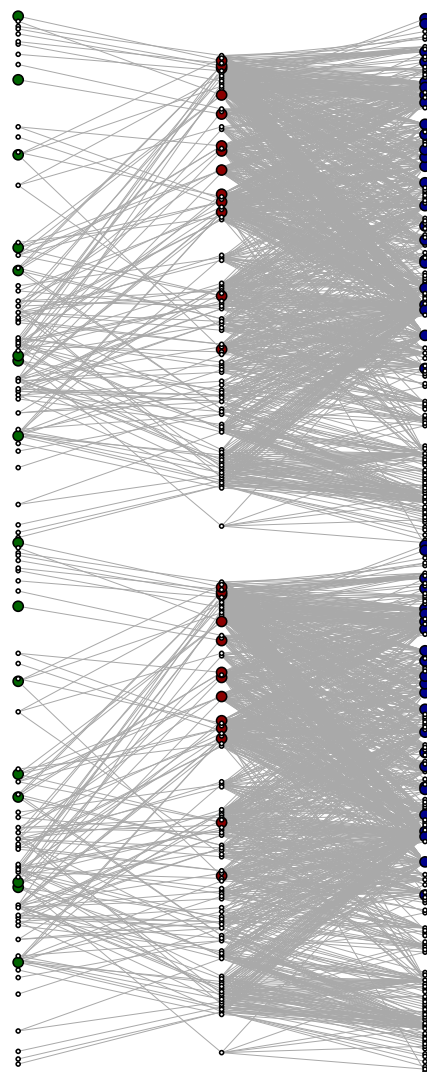
Probabilistic representation of the interaction probability between a leaf galler (*Phyllocolpa prussica*) and a parasitoid (*Chrysocharis elongata*) across Europe. The maps are generated from predicted probabilities according to the model illustrated at Fig. 3.

Figure 5


Representation of the Salix-galler and galler-parasitoid metawebs. Black cells indicate species pairs for which at least one interaction was recorded, white cells indicate absence of recorded interactions and the red cells show pairs of species never detected at the same site (and hence species pairs for which we have no information on whether they would interact should they co-occur).

Figure 6


Mapping the distribution of species richness, the number of links and connectance across Europe. The representation is based on predictions from model 3 (see Table 2). Species richness is obtained by the summation of individual occurrence probabilities and link density is obtained by the summation of the interaction probabilities.



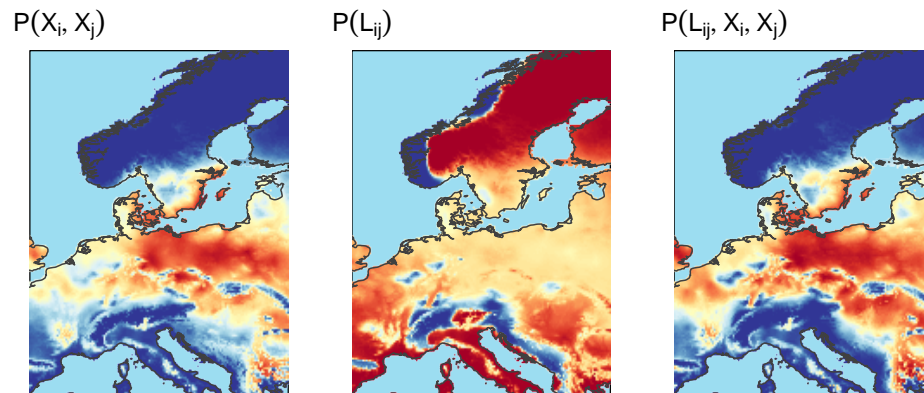
680 **Figure 2**




681 **Figure 3**



682 **Figure 4**





684 **Figure 6**

